

Final Internship Report: Design and Implementation of a Next-Generation DevOps Pipeline Managed Service on a High-Performance Cloud Platform

Name : Yash Gupta

College : Shri Ram Murti Smarak College of Engineering & Technology, Bareilly

Branch : B.Tech Computer Science

Duration : (7th May 2025 - 7th August 2025)

Email : yashg5577@gmail.com

1. Project Description

This report details the successful design, implementation, and management of a comprehensive, end-to-end automated software delivery pipeline, undertaken as part of an internship program. The project addresses the critical industry challenge of overcoming the inefficiencies, high error rates, and slow time-to-market inherent in traditional, manual software development and deployment methodologies.¹ The primary objective was to construct a state-of-the-art managed service for a web application, simulating a realistic enterprise environment where agility, reliability, and speed are paramount.

The implemented solution is a fully automated Continuous Integration and Continuous Delivery (CI/CD) pipeline built on a modern, cloud-native technology stack. The pipeline leverages Amazon Web Services (AWS) as the foundational cloud platform, providing scalable and resilient infrastructure. Infrastructure as Code (IaC) principles were adopted using Terraform, enabling the repeatable, version-controlled, and consistent provisioning of all cloud resources. Configuration management was handled by Ansible, ensuring that all server environments were configured identically and automatically. The core application was containerized using Docker, which guarantees application portability and environmental consistency from development through to production. This containerized application was then deployed and managed by Kubernetes, an industry-leading container orchestration platform that provides automated scaling, self-healing, and robust runtime management. The entire workflow, from code commit to final deployment, was orchestrated by Jenkins, which served as the central automation engine for the pipeline.¹

Key achievements of this project include the complete automation of the Software Development Lifecycle (SDLC), which has drastically reduced the need for manual intervention and minimized the potential for human error. The establishment of an IaC-driven process ensures that infrastructure is both auditable and easily reproducible across multiple environments. Furthermore, the use of containerization has solved the perennial "it works on my machine" problem, ensuring that the application behaves predictably regardless of where it is deployed.

In conclusion, this project serves as a powerful and practical demonstration of modern DevOps

principles. The resulting pipeline provides a framework for increased development velocity, significantly improved code quality, and a highly scalable and resilient application delivery system. The hands-on experience gained through this comprehensive project has provided invaluable insights into the complexities of modern IT infrastructure and has built a strong foundation for a professional career in the fields of Cloud Engineering, DevOps, and Site Reliability Engineering.¹

2. Tools & Technology Used

The successful implementation of the automated pipeline relies on a carefully selected suite of industry-leading tools and technologies. Each component in this stack has a specific and well-defined responsibility, working in concert to create a seamless and efficient workflow. This section provides a detailed examination of each technology, outlining its core concepts and its specific application within the context of this project.

2.1 Cloud Platform: Amazon Web Services (AWS)

AWS served as the foundational IaaS provider, offering the on-demand, scalable, and secure cloud infrastructure required to host the entire CI/CD pipeline and the deployed application.¹ The project leveraged a specific set of AWS services to build a robust and isolated environment:

- **Elastic Compute Cloud (EC2):** This service provided the core compute capacity in the form of virtual servers, or "instances." EC2 instances were provisioned to serve as the Jenkins controller and agent nodes, as well as the control plane and worker nodes for the Kubernetes cluster.¹
- **Virtual Private Cloud (VPC):** A VPC was established to create a logically isolated virtual network within the AWS cloud. This provided complete control over the networking environment, including the ability to define private and public subnets, configure route tables, and set up internet gateways, thereby enhancing security and network segmentation.¹
- **Simple Storage Service (S3):** S3 was utilized for its highly scalable and durable object storage capabilities. Its primary roles in the project included storing build artifacts, application logs, and, critically, the Terraform remote state file, which is essential for collaborative infrastructure management.¹
- **Elastic Load Balancing (ELB):** An Application Load Balancer (ALB), a type of ELB, was deployed to automatically distribute incoming application traffic across multiple container instances (Pods). This ensures high availability, fault tolerance, and seamless scaling by routing traffic only to healthy targets.¹

- **Identity and Access Management (IAM):** IAM was used to securely manage access to all AWS services and resources. By creating users, groups, and roles with specific policies, the project adhered to the security best practice of granting least privilege, ensuring that each component and user had only the permissions necessary to perform their function.¹
- **CloudWatch:** This service provided essential monitoring and observability for the AWS infrastructure. CloudWatch was configured to collect metrics such as EC2 CPU utilization, disk space, and network traffic, and to set up alarms that could trigger automated actions or notify administrators of potential issues.¹

2.2 Infrastructure as Code (IaC): Terraform

Terraform, by HashiCorp, was the chosen tool for implementing Infrastructure as Code. Its role was to automate the provisioning and lifecycle management of the AWS infrastructure in a declarative, predictable, and version-controlled manner.¹

- **Key Concepts:** Terraform operates on a declarative model; users define the desired *end state* of the infrastructure in configuration files written in HCL, and Terraform determines the necessary actions (API calls) to achieve that state. It uses a system of "providers" to interact with various cloud platforms, making it cloud-agnostic. A critical component is the "state file," which keeps a record of the managed infrastructure. For this project, the state file was stored remotely in an S3 bucket with state locking enabled via DynamoDB, a best practice that prevents conflicts in a multi-user environment.¹
- **Project Application:** Terraform was used to codify the entire cloud environment, including the VPC, public and private subnets, security groups, IAM roles, and the EC2 instances for the Jenkins and Kubernetes clusters. This allowed the entire infrastructure to be created, modified, or destroyed with a few simple commands (terraform apply, terraform destroy), ensuring consistency and repeatability.¹

2.3 Configuration Management: Ansible

While Terraform excels at provisioning infrastructure, Ansible was used for the subsequent step of configuration management. Its role was to automate the installation of software and the configuration of the operating systems on the EC2 instances provisioned by Terraform.¹

- **Key Concepts:** Ansible's key features include its agentless architecture, connecting to managed nodes via standard SSH, which simplifies setup. It uses a simple, human-readable YAML syntax for its "playbooks." A core principle of Ansible is idempotence, meaning that running a playbook multiple times will result in the same system state, as it only applies changes if the desired state is not already met.¹
- **Project Application:** Ansible playbooks were developed to perform tasks such as updating system packages, installing the Docker engine, and installing the necessary Kubernetes components (kubeadm, kubelet, kubectl) on each server designated to be part of the Kubernetes cluster. This automated the transformation of a bare OS into a fully functional cluster node.¹

2.4 Source Code Management & Version Control: Git and GitHub

Git and GitHub formed the foundation of the development workflow and the starting point of the CI/CD pipeline.

- **Role:** Git was used as the distributed version control system (DVCS), allowing developers to track changes to the source code locally and work on different features in parallel using branches. GitHub served as the centralized, cloud-hosted platform for the Git repositories. It facilitated collaboration among team members and, through its webhook functionality, acted as the primary trigger for the Jenkins pipeline upon every code push.¹

2.5 CI/CD Orchestration Engine: Jenkins

Jenkins was the heart of the CI/CD pipeline, acting as the central automation server that orchestrated the entire workflow from code commit to deployment.¹

- **Project Application:** A Jenkins pipeline, defined as code in a Jenkinsfile, was configured to continuously monitor the main branch of the GitHub repository. When a change was detected, Jenkins would automatically execute the defined stages: checking out the code, running Pytest for automated testing, invoking Docker to build and tag a new container image, pushing the image to Docker Hub, and finally, triggering the deployment to the Kubernetes cluster by applying the relevant manifest files.¹

2.6 Application Containerization: Docker

Docker was the chosen technology for containerizing the Python Flask application, a critical step for achieving environmental consistency and portability.¹

- **Key Concepts:** Docker packages an application and all its dependencies (libraries, runtime, system tools) into a standardized, isolated unit called a container. This is achieved by creating a "Docker image" from a set of instructions in a Dockerfile. The resulting container can run on any machine with a Docker runtime, eliminating the common "it works on my machine" problem by ensuring the application's environment is identical everywhere.¹
- **Project Application:** A Dockerfile was created to define the steps for building the Flask application image. This image was then used to create consistent and isolated environments for testing and production deployment.

2.7 Container Orchestration: Kubernetes (K8s)

For managing the containerized application in a production-like environment, Kubernetes was employed. It provides a robust platform for automating the deployment, scaling, and operation of application containers across a cluster of machines.¹

- **Key Concepts:** Kubernetes abstracts the underlying infrastructure into a unified cluster. Its fundamental components include **Pods** (the smallest deployable unit, typically holding one container), **Deployments** (which manage the desired number of replica Pods and handle rolling updates), and **Services** (which provide a stable network endpoint to access a group of Pods). The cluster is managed by a "control plane" that ensures the actual state of the cluster matches the desired state defined in YAML manifest files.¹
- **Project Application:** Kubernetes was used as the target deployment environment. The Jenkins pipeline deployed the Dockerized Flask application to the Kubernetes cluster. Kubernetes then handled the runtime management, ensuring the application was always running, automatically restarting failed containers (self-healing), and enabling easy scaling of the application by simply adjusting the number of replicas in the Deployment manifest.

2.8 Application Framework and Development Environment

The project was built around a specific application and a set of standard development tools.

- **Application:** The web application was developed using **Python** with the **Flask** framework, a lightweight and flexible choice for building web APIs and applications.¹
- **Development Tools:** **Visual Studio Code (VSCode)** was used as the primary code editor for its rich feature set and extensibility. All command-line operations were performed using **Git Bash** or a standard Linux **Terminal**.¹

2.9 DevOps CI/CD Technologies & Cloud

DevOps is a set of practices that combines software development and IT operations to shorten the development lifecycle and provide continuous delivery with high software quality.¹ At its core, it is built on principles of automation, the creation of rapid feedback loops, and a culture of shared ownership among all stakeholders in the software delivery process. The central technical artifact of the DevOps philosophy is the CI/CD pipeline.

The CI/CD pipeline is an automated workflow that guides a software change from the developer's workstation to the end user. It is composed of three interconnected, and sometimes overlapping, stages:

- **Continuous Integration (CI):** This is the practice where developers frequently—often multiple times a day—merge their code changes into a central, shared repository like Git. Each merge automatically triggers a build process and a suite of automated tests. The primary benefit of CI is the early detection of integration errors and bugs, which are far easier and less costly to fix when caught immediately rather than weeks or months later in the development cycle.¹
- **Continuous Delivery (CD):** This practice is a logical extension of Continuous Integration. It ensures that every code change that passes the automated testing phase is automatically built and packaged into an artifact that is ready to be deployed to a production environment. The final deployment to production is typically initiated by a manual approval step, giving teams control over *when* to release. Continuous Delivery ensures that the software is always in a releasable state.¹
- **Continuous Deployment (CD):** This is the most advanced stage of the pipeline and represents full automation. In a Continuous Deployment model, every code change that successfully passes through all the preceding automated stages (build, unit tests, integration tests, etc.) is automatically deployed to the production environment without any human intervention. This practice enables organizations to release new features and fixes to users rapidly and frequently.¹

2.10 Cloud Computing: The Engine for Scalable Infrastructure

Cloud computing is the on-demand delivery of IT resources—such as servers, storage, databases, and networking—over the internet with a pay-as-you-go pricing model.¹ Instead of investing heavily in and maintaining physical data centers, organizations can access these services from a cloud provider like AWS. This model provides the agility, scalability, and global reach necessary to support modern DevOps practices. The services offered by cloud providers are typically categorized into several models of abstraction, often referred to as "Everything as a Service" (XaaS).

- **Infrastructure as a Service (IaaS):** This is the most fundamental layer, providing access to core computing, networking, and storage resources. Users have control over the operating system and deployed applications but do not manage the underlying physical hardware. In this project, AWS EC2 was used as the IaaS layer to provision the virtual servers that hosted the entire pipeline and application environment.¹
- **Platform as a Service (PaaS):** This model abstracts away the underlying infrastructure, including servers and operating systems, allowing developers to focus solely on writing and deploying their application code. AWS Elastic Beanstalk is a prime example of a PaaS offering that could serve as a deployment target for the application.¹
- **Software as a Service (SaaS):** This model provides ready-to-use software applications delivered over the internet. Users interact with the software without any concern for the underlying infrastructure or platform. GitHub, used for source code management in this project, is a classic example of a SaaS provider.¹
- **Function as a Service (FaaS):** Often associated with serverless computing, FaaS abstracts away everything except the application logic itself. Developers provide code in the form of functions, and the cloud provider automatically manages all the resources required to run them in response to specific events or triggers. AWS Lambda is a leading FaaS offering that could be integrated for event-driven automation tasks.¹

These cloud services can be deployed in various models, including Public Cloud (resources owned and operated by a provider like AWS), Private Cloud (resources used exclusively by one

organization), and Hybrid Cloud (a combination of public and private clouds).¹ For this project, a public cloud model using AWS was adopted to leverage its vast array of services and scalability.

2.11 Architectural Blueprint of the CI/CD Pipeline

The architecture of the CI/CD pipeline is designed as a logical, sequential flow of automated stages. Each stage has a distinct purpose and is managed by a specific tool, creating a cohesive and orchestrated toolchain that transforms source code into a running application. The entire process is initiated by a single developer action—a code commit—and proceeds without manual intervention until the final deployment.

A visual representation of this architecture provides a clear map of the workflow, illustrating how each component interacts within the pipeline. The following narrative walkthrough describes the journey of a code change through each stage of this blueprint.

1. **Code & Commit:** The lifecycle begins on the developer's local machine. Using a code editor like VSCode and a command-line interface such as Git Bash, the developer writes or modifies code for the Python Flask web application. Once a feature or fix is complete, the changes are committed to a local Git repository. These commits are then pushed to a designated branch on a central GitHub repository, which serves as the single source of truth for the application's codebase.¹
2. **Trigger & Integration:** The push event to the GitHub repository acts as the primary trigger for the entire pipeline. A pre-configured webhook in GitHub sends a notification to the Jenkins automation server. Upon receiving this notification, Jenkins automatically initiates a new build job, connecting to the GitHub repository and pulling the latest version of the source code into its workspace. This step marks the beginning of the Continuous Integration process.²
3. **Build & Test:** Within the Jenkins environment, the pipeline executes a series of automated build and validation steps. First, it runs a suite of automated tests, such as unit tests written with Pytest, to verify the correctness and quality of the new code. If the tests pass, Jenkins

proceeds to the build phase. Using a Dockerfile located in the source code repository, Jenkins invokes the Docker engine to build the Flask application and all its dependencies into a standardized, immutable Docker image.¹

4. **Store & Version:** The newly created Docker image is a self-contained, portable artifact. To manage and version these artifacts, the pipeline tags the image with a unique identifier, such as the Jenkins build number or the Git commit hash. This tagged image is then pushed to a central container registry, like Docker Hub. The registry acts as a versioned library of application images, from which deployment environments can pull the exact version needed.¹
5. **Provision & Configure:** This stage prepares the target runtime environment. It is a two-step process driven by Infrastructure as Code and Configuration Management tools.
 - **Terraform** reads declarative configuration files (written in HCL) that define the desired state of the cloud infrastructure. It then communicates with AWS APIs to provision all necessary resources, including the Virtual Private Cloud (VPC), subnets, security groups, and the EC2 instances that will form the Kubernetes cluster.¹
 - Once the virtual servers are running, **Ansible** takes over. Using YAML-based playbooks, Ansible connects to these new instances via SSH and executes a series of configuration tasks. These tasks include installing required software like the Docker runtime and Kubernetes components (e.g., kubelet), ensuring each machine is correctly prepared to function as a node in the cluster.¹
6. **Deploy & Orchestrate:** With the application image stored and the infrastructure ready, Jenkins triggers the final deployment stage. It uses Kubernetes manifest files (e.g., deployment.yaml, service.yaml) to instruct the Kubernetes cluster on how to run the application. The Kubernetes control plane receives these instructions, schedules the application to run on its worker nodes, and pulls the specified Docker image version from the container registry. Kubernetes then deploys the application as a set of running containers (Pods) and continuously monitors their health, automatically handling tasks like scaling, networking, and self-healing.¹
7. **Expose & Monitor:** The final step is to make the application accessible to users and to monitor its ongoing health. The Kubernetes Service object automatically configures an AWS Elastic Load Balancer (ELB) to distribute incoming internet traffic to the running

application Pods, ensuring high availability. Simultaneously, AWS CloudWatch collects metrics and logs from the underlying AWS resources (e.g., EC2 CPU utilization, network traffic), providing essential data for monitoring the infrastructure's performance and stability.¹

The following table provides a structured summary of the pipeline, delineating the responsibilities of each tool at every stage. This clear separation of concerns is a hallmark of robust system design, allowing for modularity, easier debugging, and the flexibility to substitute tools as requirements evolve.

Stage	Key Action	Primary Tool(s)	Description/Purpose
Code	Develop, commit, and push source code changes.	Git / GitHub, VSCode	Manage source code versions in a distributed manner and provide a central repository for collaboration and pipeline triggers.
Build	Compile source code and package the application into a portable artifact.	Jenkins, Docker	Automate the creation of a standardized, immutable Docker image containing the application and all its dependencies.
Test	Execute automated tests to validate code quality and	Jenkins, Pytest	Run unit and integration tests automatically upon every code change to

	functionality.		provide rapid feedback and prevent regressions.
Release	Store and version the built artifact in a central repository.	Docker Hub	Act as a version-controlled registry for Docker images, ensuring reliable and consistent access to application artifacts.
Deploy	Provision infrastructure and deploy the application to the target environment.	Terraform, Ansible, Kubernetes	Automate the creation of cloud infrastructure (Terraform), configure servers (Ansible), and orchestrate the deployment of containers (Kubernetes).
Operate	Manage the running application and monitor its health and performance.	Kubernetes, AWS ELB, AWS CloudWatch	Ensure application availability, scalability, and self-healing (Kubernetes), manage traffic (ELB), and monitor infrastructure metrics (CloudWatch).

3. Conclusion: Project Outcomes and Professional Development

This internship project culminated in the successful creation of a sophisticated, fully functional DevOps pipeline for a web application deployed within a cloud computing environment. The comprehensive scope of the project, which spanned from initial application development and infrastructure provisioning to continuous, automated deployment, provided an unparalleled opportunity to bridge the gap between theoretical knowledge and practical, real-world application. The project successfully demonstrated the transformative power of automation in modern software development, establishing an efficient, repeatable, and highly scalable deployment process.¹

Summary of Achievements:

The primary achievement was the integration of a diverse toolchain—including Git/GitHub, Jenkins, Docker, Ansible, Terraform, and Kubernetes—into a cohesive and automated system. This integration yielded several key benefits. The ability to automatically build, test, and deploy code changes with a single commit dramatically improved development velocity and enhanced overall code quality. By automating the entire SDLC, the pipeline effectively minimized the risk of human error, which is a common source of failure in manual deployment processes. Furthermore, the use of Infrastructure as Code and containerization ensured absolute consistency between development, staging, and production environments, eliminating environmental drift and making deployments more predictable and reliable. The project also underscored the critical role of cloud platforms like AWS in providing the foundational, on-demand resources necessary for building scalable and resilient applications.

Technical Skills Acquired:

The hands-on nature of this project facilitated the acquisition and mastery of a wide range of critical technical skills. Proficiencies were developed in:

- **Cloud Infrastructure Management:** Deep practical experience was gained in provisioning, configuring, and managing core AWS services, including EC2, VPC, S3, and ELB.

- **Infrastructure as Code (IaC):** Expertise was developed in using Terraform to define and manage complex cloud infrastructure through declarative code.
- **Configuration Management:** Skills were honed in using Ansible to automate server configuration and software installation in an idempotent manner.
- **Containerization and Orchestration:** A thorough understanding of Docker for creating portable application containers and Kubernetes for managing those containers at scale was achieved.
- **CI/CD Orchestration:** Proficiency was gained in designing and implementing complex automation workflows using Jenkins and pipeline-as-code with Jenkinsfile.

Professional Growth:

Beyond the technical competencies, this project was instrumental in fostering a holistic understanding of the principles and culture that underpin modern IT infrastructure and application delivery. It provided a strong, practical foundation for a future career in high-demand fields such as DevOps Engineering, Cloud Architecture, or Site Reliability Engineering (SRE). The experience of building a system from the ground up has instilled a deep appreciation for the interconnectedness of development, operations, and business objectives, showcasing the ability to not only implement technical solutions but also to understand their strategic impact.¹

4. Future Scope

The pipeline developed during this internship represents a robust and functional foundation for modern software delivery. However, in the spirit of continuous improvement—a core tenet of DevOps—there are numerous avenues for enhancement that could further mature the platform. The following recommendations are presented as a logical progression, outlining a roadmap to evolve the current pipeline into a secure, observable, resilient, and enterprise-ready system.

4.1 Integrating Security: The Shift to DevSecOps

The next logical step in the pipeline's evolution is to integrate security practices directly into the automated workflow, a practice known as DevSecOps. The goal is to "shift security left," making it an early and continuous part of the development lifecycle rather than a final, often rushed, checkpoint.¹⁴

- **Recommendations:**

- **Static Application Security Testing (SAST):** Integrate SAST tools (e.g., SonarQube, Snyk) into the Jenkins pipeline. These tools would automatically scan the source code for security vulnerabilities, such as SQL injection or cross-site scripting flaws, during the build stage, providing immediate feedback to developers.¹
- **Software Composition Analysis (SCA):** Implement SCA tools to automatically scan the application's open-source dependencies (e.g., Python libraries) for known vulnerabilities (CVEs). This prevents the introduction of insecure third-party code into the application.¹⁵
- **Container Image Scanning:** Integrate tools like Trivy or Clair to scan Docker images for vulnerabilities within the OS packages and application layers before they are pushed to the container registry. This ensures that only secure images are stored and deployed.

4.2 Advanced Monitoring and Observability: Prometheus & Grafana

While AWS CloudWatch provides excellent infrastructure-level monitoring, achieving true operational excellence requires deeper, application-centric observability.¹

- **Recommendations:**

- **Deploy Prometheus:** Implement Prometheus within the Kubernetes cluster to scrape detailed, time-series metrics from applications, Kubernetes components, and the nodes themselves. Prometheus offers a powerful query language (PromQL) for in-depth analysis of performance and behavior.¹⁶
- **Visualize with Grafana:** Use Grafana to connect to the Prometheus data source and build sophisticated, real-time monitoring dashboards. These dashboards can visualize key application performance indicators (KPIs), resource utilization within the cluster, and error rates, providing actionable insights that go far beyond basic CPU and memory metrics.¹
- **Implement Automated Alerting:** Configure Prometheus Alertmanager to send proactive alerts to communication channels like Slack or PagerDuty based on predefined thresholds and complex query conditions, enabling faster incident response.

4.3 Enhancing Resilience and Scalability

To prepare the application for production-level traffic and ensure business continuity, its resilience and scalability can be further enhanced.

- **Recommendations:**

- **Implement Auto Scaling Groups (ASG):** Configure AWS Auto Scaling Groups for the Kubernetes worker nodes. This would allow the cluster's compute capacity to dynamically scale up or down based on metrics like CPU utilization, ensuring performance during traffic spikes while optimizing costs during quiet periods.¹

- **Adopt Advanced Deployment Strategies:** Evolve from the standard rolling update strategy to more advanced deployment patterns. **Blue/Green deployments** would allow for zero-downtime releases by deploying the new version alongside the old and switching traffic over instantly. **Canary deployments** would reduce risk by gradually rolling out the new version to a small subset of users before a full release.¹
- **Explore Multi-Region Deployment:** For disaster recovery and global availability, extend the Terraform and deployment configurations to support deploying the application across multiple AWS regions. This would provide high availability in the event of a regional outage.¹

4.4 Advanced Tooling and Process Maturity

To scale the DevOps practices across a larger organization, adopting enterprise-grade tooling and more modular approaches is essential.

- **Recommendations:**

- **Centralize Ansible with AAP:** For managing automation at scale, migrate from standalone Ansible to the Ansible Automation Platform (AAP) or its predecessor, Ansible Tower. These platforms provide centralized management, role-based access control (RBAC), graphical dashboards, and a REST API for better control and integration.¹
- **Develop Reusable Terraform Modules:** Refactor the existing Terraform code into reusable modules. This promotes a "Don't Repeat Yourself" (DRY) approach to infrastructure, making the code more maintainable, consistent, and easier to share across different projects and teams.¹
- **Investigate Enterprise Kubernetes Platforms:** Explore platforms like Red Hat OpenShift, which is built on top of Kubernetes. OpenShift provides an integrated set of tools for developers, enhanced security features, and streamlined workflows that can accelerate application development and deployment in a large-scale enterprise environment.¹